URP / □□□□□

□□ GPU □□□□□□□□

□□□□□□□□

URP □□□□□ · □ PtxDec □□□ Agent □□□□

PyTorch □□ → □□□□ CUDA → PTX / SASS

□□□□ □□□□ □□□□ □□□□□

□□□□□□□□□□□□□□□□□□□□□□□□

□□□□□□□□□□□□

□□□□□□□□□□□□□□□□□□□□□□□□ CUDA

□□□□□□□□□

- □□□□□□□□□□□□□□□□□□□□□□□
- □□□□□□□□□□□□□□□□□□□□□□□
- □□ CUDA □□□□□□□□□□□□□□□□□

□□□□□□□□□□□□

- □□□□□□□□ PTX□SASS □□□□□□
- □□□□□□ CUDA kernel□□□□□□□□□□□□□□
- □□□□□□□□□□launch □□□□□□□□□

□□□□□□□□□□□□□□□□□□□□□□□□ CUDA–□□□□□□□□□□

PtxDec□□□□□□□□□□□□

Enhancing LLM to Decompile Optimized PTX to Readable CUDA for Tensor Programs · ASE 2025

① □□□□□□□□

Ansor□□□□□□

Welder□□□□□□

→

② □□□□□

CUDA □□

MinHash / LSH □□

→

③ □□□□

PTX → Rolled-PTX

LoRA / SFT → CUDA

40 □□ CUDA–PTX kernel

- □□ shape□grid / block □□□□□□
- □□□□□□□□□□□□□□□□□□□□□□

Rolled-PTX □□□

- □□□□□□□□□□□□□□□□□□
- □□□□ PTX □□□□□□ 41%□
- □□□□□□□□□□□□□□□□□□ PTX□

□□□

□□□□□□□□□

□□□□□□□□□□□□□□□□□□

□□ / □□	□□□	□□□□□	I/O □□□
Qwen2.5-Coder-7B	42.16%	34.61%	29.26%
PtxDec-7B	**97.24%**	**96.40%**	**90.89%**
□□ CUDA □□	75.07%	70.15%	59.49%
□□ PTX □□	89.05%	72.17%	52.57%

□□□□□□□□

□□□□ kernel □□PtxDec-7B / 32B □ I/O □□□□ **29.10% / 60.31%**□

□□□□□□□

□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□ SASS □□□

[illegible] Agent [illegible]

[illegible]

[illegible]	PtxDec [illegible]	[illegible]
CUDA [illegible]	[illegible]	Agent [illegible] Ops-6K [illegible] PyTorch [illegible]
[illegible]	[illegible]	[illegible]
[illegible]	PTX / Rolled-PTX	[illegible] PTX[illegible]cubin[illegible]SASS
[illegible]	[illegible]	[illegible]

> [illegible]Agent [illegible]

□□□□□□—□□—□□—□□

□□□□□□□□□□□□□□□□□□□□□□

□□□□□

model.py
shape / dtype / stride
□□□□□□□□□

→

□□□□□

Python □□
CUDA kernel / launcher
C++ □□

→

□□□□

□□ / □□□□
□□
5 □ seed □□□□

□□□□□□□□□□□□□□ → □□□□ → □□□□ ↺

□□□□□

- □□□□□□□□□□□□□□□□□□□□
- □□ PTX / cubin / SASS□□□□□□□□□

□□□□□□

- □□□□□□□□□□□□□□ 3 □□□□
- □□□□ 4 □□□□□□□□□□□□□□□
- □□ API □□□□ 3 □□□□□

□□□□□□□□□□□□□□□□□□

□□□□□□□□□□□□□□□□□□□□□□□□□□□□

① □□□□□

- □□□□□□□□□□forward □□□ state_dict□
- □□ shape / dtype / stride □□□□□□□
- □□□□□□□□□□□□□□□

② □□□□

- □□□□□□□□□□□□□□□□□□□□□□□□
- □□ 5 □□□ seed □□□ forward □□□
- □□□□□□□□□□□□□□□□□□□□

③ □□□□□

- □□□□□□□□□□□□□□□□□□□□□□□□
- □□□□□□□□□□□□□□□□
- □□□□□□□□□□□□

④ □□□□

- □□□□□□□□□“□□□□□□□”□
- □□□□□□□□□□□□□□□□□□
- □□□□□□□□□□□□□ launch □□□□□

□□□□□□□□□□□□□□□

□□□□□□□□□□□ · □□□□ task 0000–0049

21

□□□□□□□□□□

20

□ CUDA / PTX / cubin / SASS

18

□□□□□□□□□□□□□□

□□□□□□□□□□35 □

17 □□ / 18 □□□□□□□□□□□□□ **48.6%**□□□ 15 □□□□□□□□□□□□□□□□□ □□□□□□□□□□

□□□□□□□□□□□□□□□□

□□□□□□□□□□□□□□□□□□□□□□□□18 □□□□□□□□□□□□ 17 □□□□□□□□□□□ □□□

18 □□□□□□□□□□□□□□□□□ kernel □□□□□□□□□□□ 32/50 □□□□□□□□□□

[illegible]

[illegible]18 [illegible]

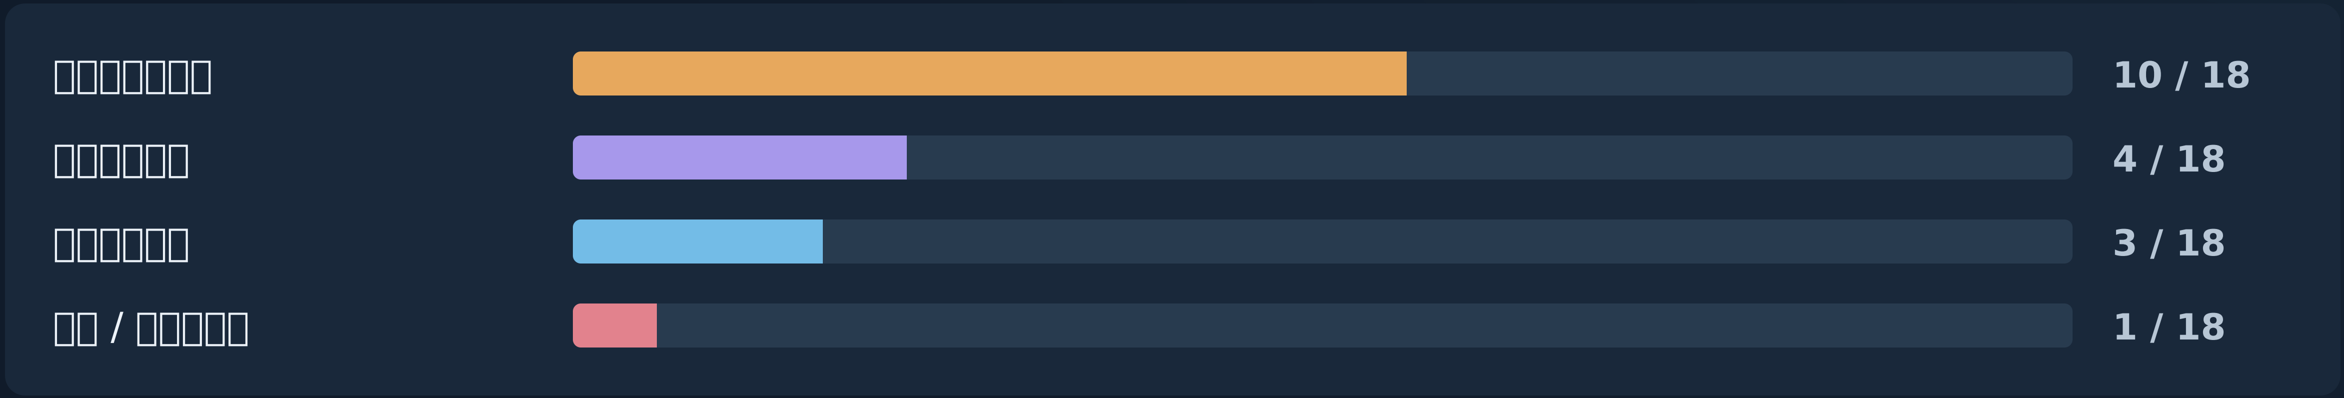

[illegible]

[illegible]shape [illegible]

[illegible]

task 31 / 36 [illegible] KeyboardInterrupt[illegible]32 / 38 [illegible]“[illegible]”[illegible]

[illegible]

[illegible]task 0,1,2,3,8,10,13,14,19,20,21,24,27,29,31,32,36,38

□□□□□Task 0006 □ ge kernel

ReflectionPad2d → ge → remainder · □□□□□□ CUDA □ SASS □□

CUDA · □□□□□□□□□□□□

```
__global__ void ge_scalar_kernel(
    const float* __restrict__ input,
    unsigned char* __restrict__ output,
    float scalar,
    int total_elements,
    int in_stride_n, int in_stride_c, int in_stride_h, int
in_stride_w,
    int H, int W, int C) {
    int idx = blockIdx.x * blockDim.x + threadIdx.x;
    int stride = blockDim.x * gridDim.x;
    for (int i = idx; i < total_elements; i += stride) {
        int w = i % W;
        int tmp = i / W;
        int h = tmp % H;
        tmp /= H;
        int c = tmp % C;
        int n = tmp / C;
        const float* in_ptr = input + n * in_stride_n + c *
in_stride_c + h * in_stride_h + w * in_stride_w;
        float val = *in_ptr;
        output[i] = (val >= scalar) ? 1 : 0;
    }
}
```

SASS · cuobjdump □□□□ / sm_89

```
/*0000*/                   IMAD.MOV.U32 R1, RZ, RZ, c[0x0][0x28] ;
/*0010*/                   S2R R3, SR_CTAID.X ;
/*0020*/                   S2R R0, SR_TID.X ;
/*0030*/                   IMAD R3, R3, c[0x0][0x0], R0 ;
/*0040*/                   ISETP.GE.AND P0, PT, R3, c[0x0][0x174],
PT ;
/*0050*/               @P0 EXIT ;
/*0060*/                   IMAD.MOV.U32 R0, RZ, RZ, c[0x0][0x0] ;
/*0070*/                   ULDC.64 UR4, c[0x0][0x118] ;
/*0080*/                   BSSY B0, 0x8f0 ;
/*0090*/                   IMAD R0, R0, c[0x0][0xc], RZ ;
/*00a0*/                   I2F.U32.RP R6, R0 ;
/*00b0*/                   IMAD.MOV R9, RZ, RZ, -R0 ;
```

□□ custom_ops.cu □□□□□ kernel□□□□□□□□□□□□□□□□“□□ CUDA □□□□ SASS”□

□□□□□□□□0002
□□□□□□□□□□□□□□□□□ kernel□

results/task_0006/kernels/custom_ops.cu:51□custom_ops.sass:664□verification.log

□□□□□□□□□□□□□□□□□□□□

□□□□□□□□□□□□□□□□□□

□□□□□

□ kernel □□
Python □□□□□□□

→

kernel □□□

□□□□□□□□□
CUDA ↔ PTX ↔ SASS

→

□□□□□□□

□□ / ABI / launch □□
□□□□□□□□□

□□□□□□□

- □□□□ / □□□□□□□□□□□□□□□
- □□□□ kernel □□□□□□ .entry / Function□
- □□□□□□□□□□□□□□□□□□

□□□□□□□□□

- □□□□□□□□□□□□□□□□□□
- □□ kernel □□□□□□□□□□□□□□□□□
- AST / token □□□□□□□□□□□□□□ 0.9□

□□□□□CUDA □□□ nvcc -ptx □ nvcc -cubin□□□ cubin □□ SASS□

□□□□□□□□ PTX □□□□□ SASS

□□□□□□□□□□□□□□□□□□

□□	□□□□	□□□□
□□□□□□□□□□□	□□ kernel□□□□□□□□□□□□ □□□□ PTX / SASS	□□□□□□□□I/O □□□
□□□□□□□□□□	□□□□□ vs □□□ / launch □□□ □□□□□□□□□□□□	□□□□□□□□□□
Agent □□□□□□□	□□□□□□□□□ □ / □□□□□□□□□□	□□□□□□□□ token □□
□□□□□□□□□□□	□□ I/O □□□□□□□□□□	□□□□□□□□□□□□□

> □ zero-shot □□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□

□□□□□□□□□

□□□□□□□□□□□□□□

P0 **□□ 18 □□□□□□□ kernel □□□**

□□

P0 **□□ PTX / SASS zero-shot □□**

□□

P1 **□□□□□□□□□□□□**

□□□□□□□□□□□□□□ 200+ □□□□□□□□□□□□□ kernel □□□

P2 **□□□□□□□□□□□□□□**

□□

□□□□□□□□□□□□□□□**PTX** □□□□□□□□□□□□□□□□ SASS□

□□□□□□□□□□□□

□□□□□□□□□

□□□□□□

- □□□results/task_* □ 21 □□□□
- □□□□□□□□□□ .cu / .ptx / .cubin / .sass□
- □□□verification.log □□ [PASS] verify success□
- □□□□□□□□□□□□□□□□□□□□□□□□□□□□

□□□□□□□□□

- □□□□ 18 □ ≠ □□□□ 17 □□
- □□□□□□□□ 15 □□□□□□□□□
- task 9 □□□□□□□□□ PTX / cubin / SASS□
- task 0 / 2 □□□□□□□□□□□□□□□□

□□□Sun et al., ASE 2025□DOI 10.1109/ASE63991.2025.00185
□□□□□□docs/urp_report/evidence.json□□□□□.md□results/□logs/□
□□□□□□□□□□□□□□□□□□□□□□□ GPU / □□□□□